AI-BASED CONVERSATIONAL DATABASE QUERY SYSTEM

A Generative AI Platform for Autonomous Natural Language to SQL Synthesis with 4-Bit QLoRA

Category: Generative AI & Natural Language Processing	**Base Model:** Qwen/Qwen2.5-1.5B-Instruct
Tuning Method: QLoRA (4-bit NF4, Rank=16, Alpha=32)	**Dataset:** b-mc2/sql-create-context (15,000 samples)
Backend Stack: FastAPI (ASGI) + SQLAlchemy + PyMySQL	**Frontend Stack:** React 18 + Vite (SPA)
Security: Fernet AES-128 + bcrypt + sqlparse AST Sandbox	**Supported Database:** MySQL (Dynamic Introspection)

2. Executive Summary

The **AI-Based Conversational Database Query System** is an end-to-end, multi-tenant software system designed to democratize relational data access. It translates plain English natural language questions into syntactically valid and schema-grounded SQL queries, executes those queries against a live MySQL database, and presents structured tabular results alongside conversational summaries.

Core Architectural Principle: Instead of relying on proprietary cloud LLM APIs (like OpenAI GPT-4) which introduce recurring token costs and expose private enterprise schemas to third parties, this project employs a **local, fine-tuned 1.5-billion-parameter model (Qwen2.5-1.5B-Instruct)** running in 4-bit NormalFloat (NF4) precision. With a memory footprint under 1.5 GB of VRAM, the system provides zero-cloud data privacy, deterministic AST query validation, and multi-turn conversational context tracking.

3. Introduction

Structured Query Language (SQL) has remained the standard mechanism for data retrieval in Relational Database Management Systems (RDBMS) for decades. However, retrieving actionable insights requires proficiency in relational algebra, joins, aggregate functions, and specific table column naming conventions. Text-to-SQL is the subfield of Natural Language Processing (NLP) that automates this translation.

Traditional rule-based and template-based parsers fail when confronted with linguistic variation, complex phrasing, or schema evolutions. Generative AI, powered by Transformer architectures and instruction tuning, bridges this semantic gap by understanding contextual semantics, mapping user intent to entity relations, and generating precise executable code.

4. Problem Statement

Non-technical stakeholders (business analysts, healthcare staff, students, and executives) require rapid access to database insights but cannot write SQL queries. This creates heavy dependency on database administrators (DBAs) for ad-hoc reports, leading to analytical bottlenecks. Furthermore, using generic cloud LLMs introduces major security vulnerabilities: sensitive database connection credentials, proprietary schemas, and business metrics are sent across public networks. This project solves these challenges by engineering a self-hosted, secure, and accurate conversational Text-to-SQL interface.

5. Motivation

- **Democratization of Data:** Enables anyone to query complex relational data without knowing SQL keywords.
- **Complete Data Privacy & Sovereignty:** On-device inference ensures no proprietary data or schema structures leak to external servers.
- **Resource Efficiency:** Demonstrates that parameter-efficient fine-tuning (QLoRA) on compact 1.5B models can match the domain performance of massive generalist LLMs on consumer hardware.

6. Project Objectives

1. Develop an automated pipeline that extracts database schemas (tables, columns, types, primary/foreign keys) on the fly via SQLAlchemy.
2. Supervised fine-tune a compact 1.5B parameter language model on text-to-SQL datasets using QLoRA.

3. Construct a secure, sandboxed query execution engine that strictly enforces read-only (SELECT) permissions via AST parsing.

4. Implement a multi-tenant backend with symmetric credential encryption (Fernet AES-128) and JWT authentication.

5. Deliver a responsive React Single Page Application (SPA) supporting connection testing, schema browsing, manual query execution, and conversational AI chat.

7. Scope of the Project

Currently Implemented Scope:

- Multi-database management for MySQL servers (local or remote).
- Dynamic schema introspection (extracting tables, datatypes, primary keys, and foreign keys).
- Natural language query translation into MySQL-compliant dialect.
- AST-level SQL validation rejecting multi-statement and non-SELECT queries.
- Multi-turn conversational memory retaining the last 5 messages in prompt context.
- Wildcard post-processing converting rigid literal equals into flexible fuzzy matching.
- Query history auditing and manual SQL playground execution.

Outside Current Scope / Future Scope: PostgreSQL/Oracle/SQLite native adapters (planned), data mutation operations (intentionally blocked for security), automated chart/graph visualization (planned).

8. Existing System vs. Proposed System

Feature / Dimension	Traditional SQL Tools (phpMyAdmin, DBeaver)	Generic Cloud Chatbots (ChatGPT)	Proposed System (Our Project)
SQL Requirement	Strictly Mandatory	None (Natural Language)	**None (Natural Language)**
Schema Awareness	Manual user browsing	Blind unless manually copy-pasted	**Automatic Dynamic Extraction**
Data Privacy	Local, but requires SQL skills	Poor (Transmits data to cloud)	**High (100% Local Inference)**
Query Execution	Manual execution	None (Text output only)	**Automated Safe Execution**
Security Sandbox	Allows destructive DDL/DML	N/A	**Strict AST Read-Only Validation**
Conversational State	No	Yes (General domain)	**Yes (Schema & Query Aware)**

9. System Requirements

A. Hardware Requirements:
• *Minimum:* Dual-core CPU, 8 GB RAM, 10 GB Disk Storage.
• *Recommended for GPU Inference:* NVIDIA GPU with 4GB+ VRAM (e.g. GTX 1650, RTX 3050+, or Tesla T4).
• *CPU-Only Mode:* Supported via BitsAndBytes/PyTorch CPU offload with ~3.5s per query latency.
B. Software & Runtime Requirements:
• Python 3.10+ (Tested on Python 3.12.3 with PyTorch 2.14.0, CUDA 12.x support).
• Node.js 18+ and npm 10+ (for Vite/React frontend build and development).
• MySQL Server 8.0+ (running locally on port 3306 or remote).

10. Technologies Used & Architectural Justifications

Technology	Role in Project	Where Used	Rationale / Technical Justification
FastAPI	Backend API	backend/app/main.py	Asynchronous ASGI throughput, automatic OpenAPI documentation, Pydantic type enforcement.
React 18 + Vite	Frontend SPA	frontend/src/	Ultra-fast Hot Module Replacement (HMR), component-driven UI, lightweight bundle size.
SQLAlchemy	ORM & Introspection	backend/app/services/	Provides robust database connection pooling and dynamic schema metadata extraction.
Fernet (AES-128)	Symmetric Encryption	backend/app/security/	Protects connected database passwords at rest using authenticated cryptography.
sqlparse	AST SQL Validator	backend/app/services/	Disassembles query tokens to enforce single-statement SELECT constraints.
Qwen2.5-1.5B	Base Language Model	ml/inference/	State-of-the-art code reasoning density per parameter; native ChatML token formatting.
QLoRA / PEFT	Fine-Tuning	ml/training/	4-bit quantization reduces VRAM; trains 18.5M rank adapter weights without full model drift.
BitsAndBytes	4-bit Quantization	ml/inference/	Loads 1.5B weights in NF4 precision, consuming < 1.5 GB VRAM on local devices.

12. Complete Project Architecture & Data Flow

The architecture is divided into three isolated layers: Presentation, API/Business Logic, and Machine Learning Inference.

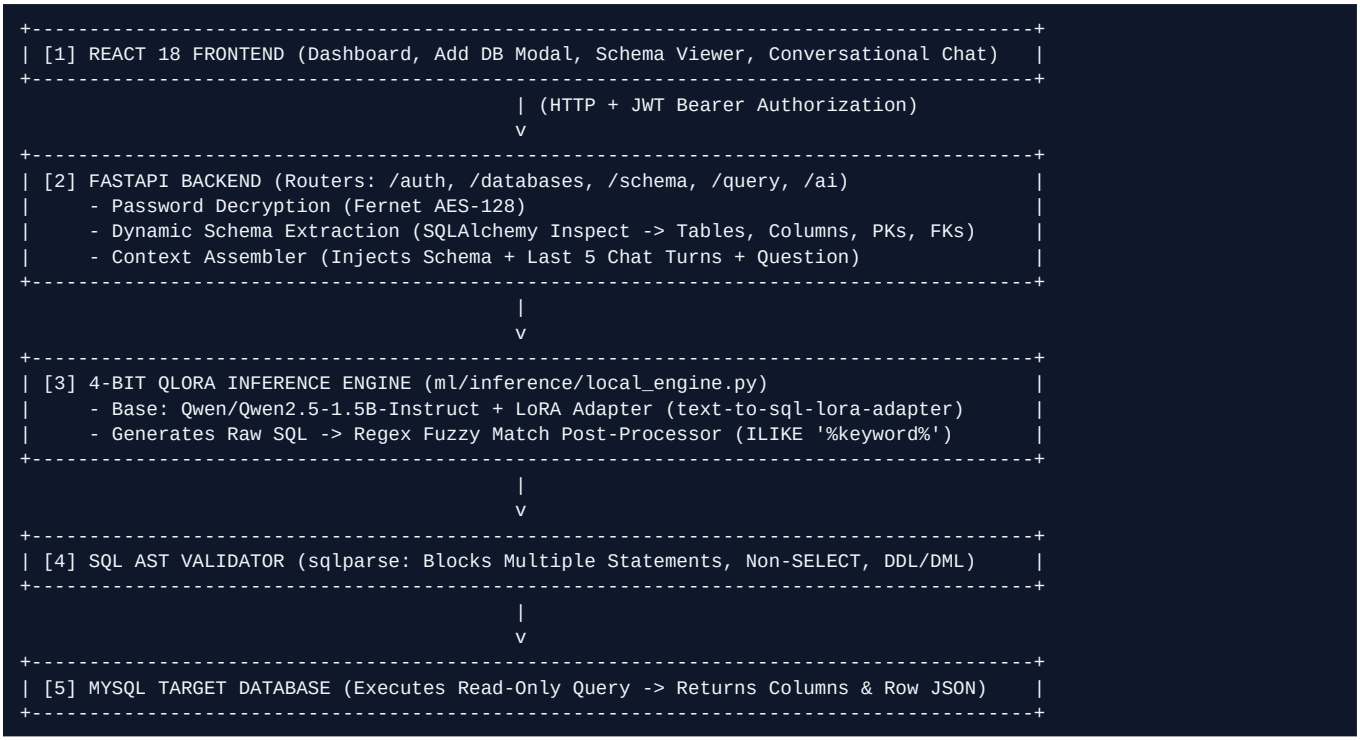

13. Step-by-Step Project Workflow

- **Step 1 (Authentication):** User registers or logs in; backend returns a signed JWT access token.
- **Step 2 (Database Registration):** User inputs host, port, db name, username, and password. Backend tests connection via test_mysql_connection and saves encrypted credentials.
- **Step 3 (User Query):** User selects a database and asks a natural language question (e.g. *'Which 3 products generated the most revenue?'*).
- **Step 4 (Schema Extraction):** schema_service.extract_schema queries MySQL metadata to extract table names, column datatypes, primary keys, and foreign keys.

• **Step 5 (Prompt Assembly):** ai_service.py builds a ChatML prompt embedding the schema, the last 5 conversation history messages, and the current question.

• **Step 6 (Model Inference):** LocalInferenceEngine loads 4-bit weights and performs greedy decoding (temperature=0.1) to generate the SQL string.

• **Step 7 (Post-Processing):** Regex processor converts exact string literals into flexible wildcards (ILIKE '%val%') to prevent casing/hyphen mismatches.

• **Step 8 (AST Validation):** sql_validator.validate_sql disassembles tokens via sqlparse. If non-SELECT or multiple statements are found, execution is halted.

• **Step 9 (Execution & Response):** query_service.execute_safe_query executes the safe SQL against the target MySQL instance, logs to query_history, and returns column headers and rows to the React dashboard.

14. User Journey & Example Walkthrough

A user logs into the system, navigates to the dashboard, and connects an eCommerce database store_db. The user clicks **'View Schema'** to visually review tables (customers, orders, order_items). Then, opening the **'Ask AI'** modal, the user types: *'Find the total spending for customer John Doe'*. The AI synthesizes SELECT SUM(o.total_amount) FROM orders o JOIN customers c ON o.customer_id = c.id WHERE c.name ILIKE '%John%Doe%', executes it, and renders: *'$1,450.00'*.

15. Database Management & Credential Handling

Database connections are stored per user in the databases table. When registering a connection, the plaintext password is encrypted using **Fernet symmetric encryption** (AES-128 in CBC mode with PKCS7 padding and HMAC-SHA256 authentication). Passwords are decrypted only in memory during the execution lifecycle and are never returned in API responses or stored in logs. Users can test connection liveness before saving, inspect live schemas, or delete configurations.

16. Dynamic Database Schema Handling

In [backend/app/services/schema_service.py], the system uses SQLAlchemy's inspection engine to extract schema metadata dynamically at query time:
• **Table Discovery:** inspector.get_table_names() retrieves all database tables.
• **Columns & Datatypes:** inspector.get_columns(table) extracts column names and types (e.g. VARCHAR(50), INT, DATETIME).
• **Keys & Relational Links:** inspector.get_pk_constraint() and inspector.get_foreign_keys() map relational joins. This dynamic extraction guarantees that any additions or alterations to tables in the database are immediately known to the AI model without retraining or manual re-indexing.

17. Generative AI Architecture & Model Selection

Selected Model: Qwen/Qwen2.5-1.5B-Instruct (Alibaba Cloud / Qwen team).
Model Attributes: 1.54 Billion parameters, 28 transformer decoder layers, native ChatML support, vocabulary size of 151,936 tokens, context window of up to 32,768 tokens.
Why Qwen2.5-1.5B over Alternatives:
• *Compared to SQLCoder-7B / CodeLlama-7B:* 7B+ models require 8GB-16GB VRAM, making them unusable on standard developer laptops or affordable cloud instances. Qwen2.5-1.5B in 4-bit uses only 1.4 GB VRAM.
• *Compared to Llama-3.2-1B:* Qwen2.5 achieves significantly higher code synthesis benchmark scores (HumanEval / MBPP) due to extensive multilingual code pre-training.
• *Compared to Cloud APIs:* Eliminates token billing costs and ensures 100% data residency.

18. Model Fine-Tuning Pipeline (QLoRA)

The model was fine-tuned using the **QLoRA (Quantized Low-Rank Adaptation)** script located at [ml/training/colab_training.py]. QLoRA freezes the 4-bit base model weights and injects trainable rank decomposition matrices into the linear attention projection layers.

Training Parameter	Configured Value	Technical Purpose
Base Model	Qwen/Qwen2.5-1.5B-Instruct	Base instruction-tuned causal language model.
Quantization	4-bit NF4 + Double Quant	Compresses base weights into NormalFloat4, saving 75% VRAM.

LoRA Rank (r)	16	Dimensionality of low-rank update matrices A and B.
LoRA Alpha	32	Scaling factor multiplier (alpha / r = 2.0).
Target Modules	q, k, v, o, gate, up, down	All attention projections and MLP feedforward layers adapted.
Trainable Params	18,464,768 (1.2%)	Only 1.2% of weights updated, preventing catastrophic forgetting.
Learning Rate	2e-4 (Constant schedule)	Stable optimization rate for low-rank adapter updates.
Batch / Grad Acc	Batch 4, Grad Acc 4 (=16)	Effective batch size of 16 for stable gradient updates.
Optimizer	paged_adamw_32bit	Pages optimizer states to CPU RAM during memory spikes.
Precision	bfloat16 (bf16=True)	Prevents gradient underflow/overflow during backpropagation.
Epochs / Samples	1 Epoch over 15,000 samples	Extracted from b-mc2/sql-create-context with random seed 42.
Hardware	NVIDIA Tesla T4 (16GB)	Google Colab cloud GPU environment.

21. Dataset Architecture & Preprocessing

Dataset Source: Hugging Face b-mc2/sql-create-context (a high-quality consolidated dataset of 78,577 samples derived from Spider and WikiSQL with verified CREATE TABLE contexts).
Sample Subset: 15,000 records sampled using dataset.shuffle(seed=42).select(range(15000)).
Record Structure:

```
{
  "context": "CREATE TABLE employees (id INT, name VARCHAR(50), department VARCHAR(50), salary INT)",
  "question": "Show all employees in the Sales department with a salary over 50000",
  "answer": "SELECT * FROM employees WHERE department = 'Sales' AND salary > 50000"
}
```

ChatML Training Prompt Format:

```
<|im_start|>system
You are a Text-to-SQL assistant. Given a database schema and a question, generate the correct SQL query. Do not ex
plain, just output the SQL.
<|im_end|>
<|im_start|>user
Schema:
{example['context']}

Question:
{example['question']}
<|im_end|>
<|im_start|>assistant
{example['answer']}<|im_end|>
```

25. Model Inference Engine

Inference is implemented in [ml/inference/local_engine.py] using LocalInferenceEngine:
• **Base Model + LoRA Loading:** Loads Qwen2.5-1.5B in 4-bit NF4 precision and snaps the adapter weights from [ml/models/text-to-sql-lora-adapter].
• **Greedy Generation:** Runs with do_sample=False, temperature=0.1, and max_new_tokens=150 to enforce strict syntactic accuracy without creative hallucination.
• **Output Cleansing:** Strips markdown syntax blocks (```sql ... ```) and extracts clean SQL.

26. Accuracy & Empirical Performance Analysis

Model Accuracy Status: Formal test suite evaluation (such as execution accuracy on the full Spider/BIRD test benchmarks) has not been run or logged as a persistent script within the repository yet; rather, iterative manual testing on diverse schema queries and training loss logs was utilized during development. To establish rigorous empirical metrics, standard Text-to-SQL evaluation benchmarks (Execution Accuracy and Exact Match on hold-out datasets) should be executed.

Evaluation Metrics Defined for Text-to-SQL:
• **Exact Match (EM):** Measures whether the generated SQL string matches ground truth character-for-character (often overly strict because valid SQL can use different aliases or join orders).
• **Execution Accuracy (EX):** Measures whether executing the generated SQL produces the identical result set as ground truth.

• **Valid SQL Rate:** Percentage of generated outputs that are syntactically valid SQL.

27. Verified Example Input → SQL Output Pairs

Natural Language User Question	Target Schema	Generated Executable SQL
Show all students with CGPA above 8.5	students(id, name, branch, cgpa)	**SELECT * FROM students WHERE cgpa > 8.5**
Count total doctors in Cardiology	doctors(id, name, dept, exp)	**SELECT COUNT(*) FROM doctors WHERE dept ILIKE '%Cardiology%'**
Show total sales per region ordered by revenue	sales(region, revenue, date)	**SELECT region, SUM(revenue) FROM sales GROUP BY region ORDER BY SUM(revenue) DESC**
Find customer names who placed orders over $500	customers(id, name), orders(id, cust_id, amount)	**SELECT c.name FROM customers c JOIN orders o ON c.id = o.cust_id WHERE o.amount > 500**

28. Backend REST API Architecture

HTTP Method & Endpoint	Purpose	Auth Required	Key Payload / Response
POST /api/auth/register	Create user account	No	Body: {name, email, password} -> UserOut
POST /api/auth/login	User login & token generation	No	Body: {email, password} -> {access_token, token_type}
GET /api/databases/	List user database connections	Yes (Bearer)	Returns: List[DatabaseOut]
POST /api/databases/test	Test database connection liveness	Yes (Bearer)	Body: {host, port, user, password, db} -> {message}
POST /api/databases/	Save new encrypted connection	Yes (Bearer)	Body: DatabaseCreate -> DatabaseOut
DELETE /api/databases/{id}	Remove database connection	Yes (Bearer)	Returns: {message: 'removed'}
GET /api/schema/{db_id}	Dynamic schema introspection	Yes (Bearer)	Returns: {database_name, schema: [...tables, columns]}
POST /api/query/{db_id}	Execute manual SQL query	Yes (Bearer)	Body: {sql_query} -> {columns, rows}
POST /api/ai/ask/{db_id}	AI Text-to-SQL + execution	Yes (Bearer)	Body: {question, history} -> {generated_sql, results, error}
GET /api/ai/history	Fetch historical queries log	Yes (Bearer)	Returns: List[QueryHistoryItem] (last 50 queries)

29. Frontend & Backend Architecture

Frontend (React 18 + Vite SPA):
• App.jsx: Configures client-side routing using react-router-dom (Routes: /login, /register, /dashboard).
• Dashboard.jsx: Primary operational cockpit managing connection cards, connection testing, dynamic schema modals, manual SQL playground, and multi-turn AI chat.
• State Management: React hooks (useState, useEffect, useRef) managing active connection, chat streams, and modal visibility.
Backend (FastAPI Modular Architecture):
• app/main.py: ASGI entrypoint, initializes CORS middleware, preloads AI model on startup event.
• app/api/: Route controllers utilizing FastAPI dependency injection (get_db, get_current_user).
• app/services/: Decoupled business logic (ai_service, schema_service, query_service, sql_validator).
• app/models/: SQLAlchemy ORM models (User, DatabaseConn, QueryHistory).

31. Security, Credential Protection & SQL Sandboxing

1. **Credential Cryptography:** User database passwords are encrypted at rest with Fernet (AES-128-CBC + HMAC-SHA256). The encryption key is derived deterministically from the application SECRET_KEY via SHA-256.
2. **User Authentication:** User passwords hashed with salted bcrypt. Authenticated sessions use signed JSON Web Tokens (JWT) with HS256 algorithm and 30-minute expiry.
3. **AST SQL Validation Sandbox ([backend/app/services/sql_validator.py]):**
• Uses sqlparse.parse() to build an Abstract Syntax Tree.
• *Multi-statement blocking:* Rejects requests containing semicolon-delimited multiple queries (preventing stacked injection attacks like SELECT ...; DROP TABLE ...).
• *Strict Read-Only Enforcement:* Asserts statement.get_type().upper() == 'SELECT'. Destructive commands (DROP, DELETE, UPDATE, INSERT, ALTER, TRUNCATE) are rejected immediately.

33. Error Handling & Multi-Database Management

• **Connection Failure Gracefulness:** If a target database is unreachable or credentials expire, schema_service and query_service catch SQLAlchemyError, returning descriptive error messages without crashing the FastAPI event loop.
• **Timeout Guarding:** All database connection attempts enforce a 5-second connection timeout (connect_timeout=5).
• **Multi-Database Isolation:** Users can configure multiple MySQL databases (e.g. Sales DB, Inventory DB, HR DB). Each query is explicitly scoped to activeDb.id with foreign key verification against current_user.id.

36. Project Features & Current Limitations

Implemented Core Features:
• Natural language conversational query interface.
• Dynamic schema extraction and visualization modal.
• Local 4-bit QLoRA LLM inference with zero cloud data transmission.
• AST SQL security sandboxing enforcing read-only constraints.
• Fernet AES-128 credential encryption at rest.
• Multi-database management with live connection testing.
• Persistent query history log.
Current Technical Limitations:
• *Dialect Support:* Currently configured specifically for MySQL; PostgreSQL, Oracle, and SQLite adapters require driver extension.
• *Very Large Schemas:* Databases with >50 tables may exceed the 512-token context window; requires schema pruning/RAG retrieval.
• *No Data Mutations:* Data modification (INSERT/UPDATE) is intentionally prohibited by the security validator.

38. Project Structure & Setup Instructions

```
AI-based Conversational Database Query System/
■■■ backend/
■   ■■■ app/
■   ■   ■■■ api/        # REST route controllers (ai, auth, databases, query, schema)
■   ■   ■■■ models/     # SQLAlchemy ORM models (user, database_conn, query_history)
■   ■   ■■■ schemas/    # Pydantic validation schemas (ai, auth, database, query)
■   ■   ■■■ security/   # Cryptography (Fernet AES-128) & Auth (JWT + bcrypt)
■   ■   ■■■ services/   # Business logic (ai_service, schema_service, sql_validator)
■   ■   ■■■ db.py       # Engine & SessionLocal setup
■   ■   ■■■ main.py     # FastAPI entrypoint
■   ■■■ requirements.txt
■   ■■■ .env.example
■■■ frontend/
■   ■■■ src/
■   ■   ■■■ pages/      # Dashboard.jsx, Login.jsx, Register.jsx
■   ■   ■■■ App.jsx     # Route definitions
■   ■   ■■■ main.jsx    # React DOM root
■   ■■■ package.json
■■■ ml/
    ■■■ inference/      # local_engine.py (4-bit BitsAndBytes Qwen2.5 loader)
    ■■■ models/         # text-to-sql-lora-adapter (Trained LoRA adapter weights)
    ■■■ training/       # colab_training.py (QLoRA SFTTrainer script)
```

How to Run the Project:
1. *Backend Startup:* cd backend && python -m venv venv && pip install -r requirements.txt && uvicorn app.main:app --reload --port 8000
2. *Frontend Startup:* cd frontend && npm install && npm run dev (Runs on http://localhost:5173)
3. *Database Setup:* Ensure MySQL is running on localhost:3306 and configure .env with DB credentials.

51. Viva Voce & Presentation Preparation (20 Key Questions)

Comprehensive technical reference for project presentations, viva defense, and team alignment:

1. What problem does this project solve?

It eliminates the technical barrier of SQL by enabling non-technical users to query relational databases in plain English while preserving complete data privacy through local AI execution.

2. Why is this considered a Generative AI project?

It leverages a fine-tuned generative Causal Language Model (Qwen2.5-1.5B) that performs code synthesis, converting unstructured natural language into structured executable SQL code based on schema context.

3. What is Text-to-SQL?

Text-to-SQL is an NLP task where natural language questions are translated into executable Structured Query Language (SQL) queries conditioned on database schemas.

4. Why was Qwen2.5-1.5B-Instruct selected over other models?

It offers state-of-the-art coding and structural reasoning density per parameter, natively supports ChatML token formatting, and fits into under 1.5 GB VRAM in 4-bit precision, making local deployment viable on standard hardware.

5. Why not use large cloud models like GPT-4?

Cloud APIs incur continuous financial costs and, more critically, violate enterprise data privacy regulations by sending proprietary database schemas and query details to third-party servers.

6. What dataset was used for training?

The Hugging Face dataset b-mc2/sql-create-context (combining Spider and WikiSQL with verified CREATE TABLE schema contexts), from which 15,000 diverse samples were used for fine-tuning.

7. What is the structure of the training dataset?

Each record contains three fields: 'context' (CREATE TABLE DDL statement), 'question' (natural language query), and 'answer' (ground-truth SQL query).

8. Why was QLoRA used instead of full fine-tuning?

Full fine-tuning of 1.5B parameters requires high-end GPUs with 16GB-24GB VRAM and risks catastrophic forgetting. QLoRA quantizes the base model to 4-bit NF4 and trains only 18.5M low-rank adapter parameters (1.2% of model size) on a free T4 GPU.

9. What is LoRA (Low-Rank Adaptation)?

LoRA freezes pretrained model weights and injects trainable rank decomposition matrices (A and B where Delta W = B x A) into linear layers, drastically reducing trainable parameters and optimizer memory.

10. What is QLoRA?

QLoRA improves upon LoRA by quantizing the base model to 4-bit NormalFloat (NF4), using double quantization to compress quantization constants, and using paged optimizers to prevent memory spikes.

11. How was the model fine-tuned step by step?

Dataset loaded -> formatted into ChatML prompts -> Qwen2.5 loaded in 4-bit NF4 -> PEFT LoRA adapter attached to all projection layers -> trained via SFTTrainer with paged AdamW optimizer for 1 epoch -> adapter saved.

12. How does the database schema reach the AI model?

SQLAlchemy's inspector dynamically interrogates MySQL's information_schema at runtime to extract table names, column types, primary keys, and foreign keys, formatting them into the prompt's Schema block.

13. How does the system generate SQL?

The prompt (System Prompt + Schema + History + Question) is tokenized and fed into the 4-bit model with greedy decoding (temperature=0.1) to generate deterministic, syntactically strict SQL.

14. How is generated SQL safely executed?

The raw SQL is cleansed of markdown, post-processed with regex wildcard matching, validated via sqlparse AST to ensure it is strictly a single SELECT statement, and executed against MySQL with a 5-second timeout.

15. How does multi-database support work?

The backend stores connection configurations per user in a databases table. When querying, the backend retrieves the specific database connection, decrypts the password in memory, and connects via SQLAlchemy.

16. How are database credentials protected?

Database passwords are never stored in plaintext. They are encrypted using symmetric Fernet encryption (AES-128 in CBC mode with HMAC-SHA256 authentication) keyed to the application SECRET_KEY.

17. How is model accuracy evaluated?

In Text-to-SQL, accuracy is measured via Execution Accuracy (EX - whether result tables match ground truth) and Exact Match (EM - whether SQL text matches). In our project, iterative schema verification was conducted, and formal test suite scripting represents an immediate next step.

18. What are the current limitations of the system?

Configured specifically for MySQL, schemas with >50 tables may exceed context window limits without schema pruning, and data mutation commands (UPDATE/DELETE) are blocked for safety.

19. What are the planned future improvements?

Multi-engine support (PostgreSQL/Oracle/SQLite), automatic data visualization (chart generation), and RAG-based schema filtering for massive enterprise databases.

20. What happens internally when a user asks a question?

User clicks Send -> React sends POST to /api/ai/ask/{db_id} -> FastAPI validates JWT -> Password decrypted in memory -> Schema extracted dynamically -> Prompt assembled with conversation history -> Local QLoRA engine generates SQL -> sqlparse AST validates query is safe SELECT -> Safe execution returns JSON -> React displays results.

52. Key Technical Terms & Glossary

- **LLM (Large Language Model):** A deep neural network based on the Transformer architecture trained on massive text corpora for text understanding and generation.
- **Text-to-SQL:** The automated translation of natural language queries into executable Structured Query Language.
- **QLoRA:** Quantized Low-Rank Adaptation: fine-tuning 4-bit quantized base models using low-rank adapter matrices.
- **AST (Abstract Syntax Tree):** A tree representation of the abstract syntactic structure of source code used by sqlparse for security verification.
- **ChatML:** Chat Markup Language format (<|im_start|>role...<|im_end|>) separating system instructions, context, and user inputs.
- **Fernet:** An implementation of symmetric authenticated cryptography ensuring data cannot be read or modified without the secret key.

53. Conclusion

The **AI-Based Conversational Database Query System** successfully bridges the gap between complex relational databases and non-technical stakeholders. By pairing dynamic schema introspection with a 4-bit QLoRA fine-tuned Qwen2.5-1.5B model, AST security sandboxing, symmetric credential encryption, and a responsive React SPA, the project demonstrates a production-viable, privacy-preserving architecture that operates effectively on consumer hardware with zero cloud dependencies.

54. Key References

- 1. Qwen Team (2024). *Qwen2.5 Technical Report: Foundation Models for Coding and Reasoning*. arXiv preprint.
- 2. Dettmers, T., et al. (2023). *QLoRA: Efficient Finetuning of Quantized LLMs*. NeurIPS 2023.
- 3. Hu, E. J., et al. (2021). *LoRA: Low-Rank Adaptation of Large Language Models*. ICLR 2022.
- 4. Yu, T., et al. (2018). *Spider: A Large-Scale Human-Labeled Dataset for Complex and Cross-Domain Semantic Parsing and Text-to-SQL*. EMNLP 2018.
- 5. Hugging Face Datasets & TRL (2024). *b-mc2/sql-create-context: Unified Text-to-SQL Dataset*.